

**Performance Assessment for
D208: Predictive Modeling
Task 2 Attempt 2**

Drew Mendez
MSDA Western Governors University
D208: Predictive Modeling
Dr. Taylor Jensen
July 8, 2024

D208_PA_MendezD_Task2_Attempt2

July 8, 2024

1 Part I: Research Question

1.1 A. Purpose of the Data Analysis

1.1.1 A1. Research Question

The data set selected for this performance assessment is the **churn** data set. The research question for this assessment is:

What variables contribute to whether or not a customer will churn?

1.1.2 A2. Goals of the Data Analysis

The goal of this data analysis is to determine influential explanatory variables to develop a logistic regression model that can be used by stakeholders to predict the likelihood of the target variable **Churn**, a binary categorical variable. This model could then be used by stakeholders to determine what services or customer attributes are associated with the probability that they will churn.

2 Part II: Method Justification

2.1 B. Multiple Logistic Regression Methods

2.1.1 B1. Four Assumptions of Logistic Regression

- independence of errors
- log odds (logit) of the dependent variable is a linear combination of the explanatory variables
- no multicollinearity
- no strongly influential outliers

(Stoltzfus, 2011)

2.1.2 B2. Benefits of using Python

Python was chosen for the functionality of the many packages available and for the ability to write functions that can be reused throughout the project. The libraries and packages that are essential to this analysis are:

- From **Pandas**, the `.isnull()`, `.duplicated()`, and `.sum()` methods provide some important basic functionality. Additionally, `.quantile()` is used in the detection of outliers and `.value_counts()` is used to count unique elements in data frames.
- From **Matplotlib**, `pyplot` is used to generate univariate graphs of variables to observe their distributions and outliers.

- From **Seaborn**, `boxplot()` and `countplot()` are used to generate the bivariate graphs.
- From **Statsmodels**, `.add_constant()` and `Logit()` are necessary to produce the logistic regression model.
- From **sklearn.metrics**, `confusion_matrix()` and `accuracy_score()` were necessary to respond to E2.
- From **scipy**, `stats.chi2.sf()` was necessary to calculate the test statistic for the Likelihood Ratio Test.

2.1.3 B3. Why Logistic Regression is Appropriate

The research question posed here has a binary categorical variable as the target variable. Since the goal of this analysis is to construct a model using multiple explanatory variables to predict the likelihood of a categorical response variable, a logistic regression model is appropriate.

3 Part III: Data Preparation

3.1 C. Summary of the Data Preparation Process

3.1.1 C1. Data Cleaning Goals

Before constructing the model, first the data set will be cleaned and treated. The plan to clean the data set involves detecting and treating duplicates, missing values, and outliers, and the re-expression of categorical variables. The steps and techniques necessary to perform these tasks is given:

- Duplicates:
 - Duplicates are detected below by chaining the `.duplicated()` and `.sum()` methods from the **Pandas** library and calling them on the data frame, returning the total count of duplicate observations. It is shown below that there are no duplicate rows.
- Missing Values:
 - Missing values are detected below by chaining the `.isnull()` and `.sum()` methods from the **Pandas** library and calling them on the data frame, returning the total count of missing values for each variable. Since the **InternetService** variable has **None** as one of its options, the 2129 erroneously identified null values are imputed with **None** to avoid being interpreted as nulls. It is shown below that there are no other missing values.
- Outliers:
 - The outliers of the following thirteen **quantitative variables** are identified using the `boxplot()` function from the **matplotlib** library and counted below using a user-defined function. It is shown below that there are no unacceptable/unreasonable outliers, so all outliers shown here will be retained.
 - * Income: Annual income of customer (**continuous numeric data**)
 - * Lat: GPS coordinates of the latitude of the customer residence (**continuous numeric data**)
 - * Lng: GPS coordinates of the longitude of the customer residence (**continuous numeric data**)
 - * Population: Population within a mile radius of customer (**discrete numeric data**)
 - * Children: Number of children in customer's household (**discrete numeric data**)
 - * Age: Age of customer (**continuous numeric data**)

- * Outage_sec_perweek: Average number of seconds per week of system outages in the customer’s neighborhood (**continuous numeric data**)
- * Email: Number of emails sent to the customer in the last year (marketing or correspondence) (**discrete numeric data**)
- * Contacts: Number of times customer contacted technical support (**discrete numeric data**)
- * Yearly_equip_failure: The number of times customer’s equipment failed and had to be reset/replaced in the past year (**discrete numeric data**)
- * Tenure: Number of months the customer has stayed with the provider (**continuous numeric data**)
- * MonthlyCharge: The amount charged, on average, per customer monthly (**continuous numeric data**)
- * Bandwidth_GB_Year: The average amount of data used, in GB, in a year by the customer (**continuous numeric data**)
- Re-expression of Categorical Variables:
 - The following thirteen **binary nominal categorical variables** are re-expressed below using binary encoding:
 - * Churn: Whether the customer discontinued service within the last month (yes, no)
 - * Techie: Whether the customer considers themselves technically inclined (yes, no)
 - * Port_modem: Whether the customer has a portable modem (yes, no)
 - * Tablet: Whether the customer owns a tablet such as iPad, Surface, etc. (yes, no)
 - * Phone: Whether the customer has a phone service (yes, no)
 - * Multiple: Whether the customer has multiple lines (yes, no)
 - * OnlineSecurity: Whether the customer has an online security add-on (yes, no)
 - * OnlineBackup: Whether the customer has an online backup add-on (yes, no)
 - * DeviceProtection: Whether the customer has device protection add-on (yes, no)
 - * TechSupport: Whether the customer has a technical support add-on (yes, no)
 - * StreamingTV: Whether the customer has streaming TV (yes, no)
 - * StreamingMovies: Whether the customer has streaming movies (yes, no)
 - * PaperlessBilling: Whether the customer has paperless billing (yes, no)
 - The following six **nominal categorical variables** are re-expressed below using one-hot encoding:
 - * Area: Area type (rural, urban, suburban)
 - * Marital: Marital status of customer
 - * Gender: Customer self-identification as male, female, or nonbinary
 - * Contract: The contract term of the customer (month-to-month, one year, two year)
 - * InternetService: Customer’s internet service provider (DSL, fiber optic, None)
 - * PaymentMethod: The customer’s payment method (electronic check, mailed check, bank (automatic bank transfer), credit card (automatic))
 - If variables have too many unique values, re-expressing them will increase the dimensionality of the model (Middleton, 2022). For this reason, the following categorical variables will be omitted from the model:
 - * City: Customer city of residence
 - * State: Customer state of residence
 - * County: Customer county of residence
 - * Zip: Customer zip code of residence
 - * TimeZone: Time zone of customer residence based on customer sign-up information
 - * Job: Job of the customer/invoiced person

```
[1]: ## C1 The following cells include the annotated code used to clean the data.
# See code attached, in D208_PA_MendezD_Task2_Attempt2.ipynb

# Import the Pandas library, then load the data into a data frame with Pandas' .
  ↳ read_csv() function
import pandas as pd
df = pd.read_csv('/Users/drewmendez/Documents/WGU/D208/churn_d208/churn_clean.
  ↳ csv')

def printDupesNulls(data_frame):
# Detect duplicates with Pandas' .duplicated method chained with .sum() method.
# Identify missing values in the data frame with Pandas' .isnull() method,
# then sum the resulting series with the .sum() method

    duplicate_count = data_frame.duplicated().sum()
    missing_values_count = data_frame.isnull().sum()
    print('Number of duplicate rows:', duplicate_count)
    print("Number of missing values per variable:")
    print(missing_values_count)

printDupesNulls(df)
```

```
Number of duplicate rows: 0
Number of missing values per variable:
CaseOrder          0
Customer_id        0
Interaction         0
UID                0
City               0
State              0
County             0
Zip                0
Lat                0
Lng                0
Population         0
Area               0
TimeZone           0
Job                0
Children           0
Age                0
Income             0
Marital            0
Gender             0
Churn              0
Outage_sec_perweek 0
Email              0
Contacts           0
```

Yearly_equip_failure	0
Techie	0
Contract	0
Port_modem	0
Tablet	0
InternetService	2129
Phone	0
Multiple	0
OnlineSecurity	0
OnlineBackup	0
DeviceProtection	0
TechSupport	0
StreamingTV	0
StreamingMovies	0
PaperlessBilling	0
PaymentMethod	0
Tenure	0
MonthlyCharge	0
Bandwidth_GB_Year	0
Item1	0
Item2	0
Item3	0
Item4	0
Item5	0
Item6	0
Item7	0
Item8	0

dtype: int64

[2]: *## C1 Treatment of NAs*

```
# Since the 'InternetService' variable has 'None' as one of its options,
# it is necessary to impute 'None'
```

```
df['InternetService'].fillna('None', inplace=True)
```

```
# Verify that 'None' no longer appears as 'Null'
```

```
print('Number of `InternetService` nulls:', df['Tenure'].isnull().sum())
```

Number of `InternetService` nulls: 0

[3]: *## C1 Detect and Count Outliers of Numeric Variables*

```
import matplotlib.pyplot as plt
```

```
def boxplotOutliers(data_frame, col_name):
```

```
# Visualize outliers using boxplot() from matplotlib
```

```
# First and third quartiles, Q1 and Q3, are found using .quantile() from Pandas,
```

```

# then the interquartile range is found using  $IQR = Q3 - Q1$ .
# The upper whisker of the boxplot is found using  $max = Q3 + 1.5 * IQR$ .
# The lower whisker of the boxplot is found using  $min = Q1 - 1.5 * IQR$ .
# The .sum() method returns the count of observations greater than the max or
↳ less than the min.
# The .round() method rounds the outlier count to two decimals.
# If loop to print corresponding outputs

plt.boxplot(data_frame[col_name])
plt.title(f'Boxplot of {col_name}')
plt.show()

Q1 = data_frame[col_name].quantile(0.25)
Q3 = data_frame[col_name].quantile(0.75)
IQR = Q3 - Q1
maximum = round(Q3 + 1.5 * IQR, 2)
minimum = round(Q1 - 1.5 * IQR, 2)
outlier_count_up = (data_frame[col_name] > maximum).sum()
outlier_count_low = (data_frame[col_name] < minimum).sum()

if outlier_count_up > 0:
    if outlier_count_low > 0:
        print(f'For the `{col_name}` variable, all observations greater
↳ than {maximum} or less than {minimum} are considered outliers.')
        print(f'The count of observations greater than {maximum} is
↳ {outlier_count_up}.')
        print(f'The count of observations less than {minimum} is
↳ {outlier_count_low}.')
    if outlier_count_low == 0:
        print(f'For the `{col_name}` variable, all observations greater
↳ than {maximum} are considered outliers.')
        print(f'The count of observations greater than {maximum} is
↳ {outlier_count_up}.')
    if outlier_count_up == 0:
        if outlier_count_low > 0:
            print(f'For the `{col_name}` variable, all observations less than
↳ {minimum} are considered outliers.')
            print(f'The count of observations less than {minimum} is
↳ {outlier_count_low}.')
        if outlier_count_low == 0:
            print(f'There are no outliers for the `{col_name}` variable.')

```


[4]: # C1 Detection of Outliers for 13 Numeric Variables

```

numericVars = df[['Lat', 'Lng', 'Population', 'Children', 'Age', 'Income', '
↳ 'Outage_sec_perweek', 'Email',
                'Contacts', 'Yearly_equip_failure', 'Tenure', 'MonthlyCharge', '
↳ 'Bandwidth_GB_Year']]

for col in numericVars:
    boxplotOutliers(df, col)


```



For the `Lat` variable, all observations greater than 52.25 or less than 25.19 are considered outliers.


The count of observations greater than 52.25 is 77.

The count of observations less than 25.19 is 81.




For the `Lng` variable, all observations less than -122.57 are considered outliers.

The count of observations less than -122.57 is 273.




For the `Population` variable, all observations greater than 31813.0 are considered outliers.

The count of observations greater than 31813.0 is 937.




For the `Children` variable, all observations greater than 7.5 are considered outliers.

The count of observations greater than 7.5 is 401.




There are no outliers for the `Age` variable.



For the `Income` variable, all observations greater than 104278.35 are considered outliers.


The count of observations greater than 104278.35 is 336.



For the `Outage_sec_perweek` variable, all observations greater than 17.9 or less than 2.09 are considered outliers.

The count of observations greater than 17.9 is 43.


The count of observations less than 2.09 is 33.



For the `Email` variable, all observations greater than 20.0 or less than 4.0 are considered outliers.


The count of observations greater than 20.0 is 15.

The count of observations less than 4.0 is 23.




For the `Contacts` variable, all observations greater than 5.0 are considered outliers.

The count of observations greater than 5.0 is 8.




For the `Yearly equip_failure` variable, all observations greater than 2.5 are considered outliers.

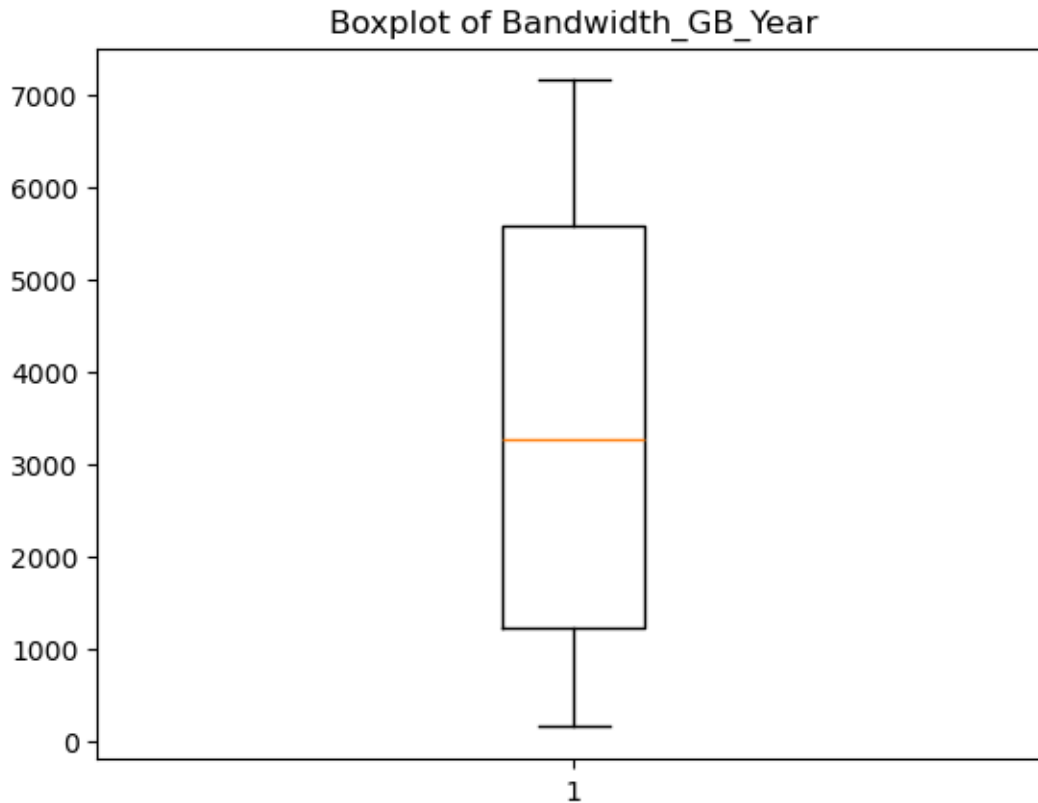
The count of observations greater than 2.5 is 94.



There are no outliers for the `Tenure` variable.



There are no outliers for the `MonthlyCharge` variable.



There are no outliers for the `Bandwidth_GB_Year` variable.

```
[5]: ## C1 Binary Encoding Re-expression of the Thirteen Binary Variables

# It was necessary to keep a copy of the Churn variable as strings for the
  ↳ bivariate graphs
df['ChurnStr'] = df['Churn'].copy()

# Create a list of the columns that will be encoded
binaryList = ['Churn', 'Techie', 'Port_modem', 'Tablet', 'Phone',
              'Multiple', 'OnlineSecurity', 'OnlineBackup',
  ↳ 'DeviceProtection',
              'TechSupport', 'StreamingTV', 'StreamingMovies',
  ↳ 'PaperlessBilling']

# Run a loop that replaces all 'Yes' with 1 and 'No' with 0 for each column in
  ↳ the list above
for col in binaryList:
    df[col] = df[col].replace({'Yes': 1, 'No': 0})

binaryVars = df[['Techie', 'Port_modem', 'Tablet', 'Phone',
```

```

        'Multiple', 'OnlineSecurity', 'OnlineBackup',
        ↪ 'DeviceProtection',
        'TechSupport', 'StreamingTV', 'StreamingMovies',
        ↪ 'PaperlessBilling']]

```

```

[6]: ## C1 One-Hot Encoding Re-expression of the Six Categorical Variables

# Data frame of categorical variables
catVars = df[['Area', 'Marital', 'Gender', 'Contract', 'InternetService',
        ↪ 'PaymentMethod']]

# Create additional data frame from variables being re-expressed
oneHotVars = catVars

# Apply one-hot encoding, dropping first column to avoid multi-collinearity in
        ↪ the model
oneHotVars = pd.get_dummies(oneHotVars, drop_first = True, dtype = int)

```

```

[7]: ## C1 Type Casting Ordinal and Binary Variables

# This was necessary for graphs below to behave correctly
ordinalVars = df[['Item1', 'Item2', 'Item3', 'Item4', 'Item5', 'Item6',
        ↪ 'Item7', 'Item8']]

binaryVars = binaryVars.astype('category')

df['Churn'] = df['Churn'].astype('category')

# assign `Churn` to depVar
depVar = df['Churn']

```

```

[8]: ## C1 Prepared Data Frame

df_prep = pd.concat([depVar, numericVars, binaryVars, oneHotVars, ordinalVars],
        ↪ axis = 1)

```

3.1.2 C2. Summary Statistics of Dependent Variable and Independent Variables

The dependent variable for this logistic regression model will be the binary categorical variable Churn.

The independent variables for the initial model will be:

- all thirteen numeric variables
 - Lat, Lng, Population, Children, Age, Income, Outage_sec_perweek, Email, Contacts, Yearly_equip_failure, Tenure, MonthlyCharge, Bandwidth_GB_Year
- all twelve re-expressed binary variables
 - Techie, Port_modem, Tablet, Phone, Multiple, OnlineSecurity,

OnlineBackup, DeviceProtection, TechSupport, StreamingTV, StreamingMovies, PaperlessBilling

- all six categorical variables re-expressed with one-hot encoding
 - Area, Marital, Gender, Contract, InternetService, PaymentMethod
- all eight ordinal variables
 - Item1, Item2, Item3, Item4, Item5, Item6, Item7, Item8

The written descriptions of the summary statistics for these variables are shown below.

[9]: *## C2 Summaries of Numeric and Ordinal Variables*

```
print('Summaries of Numeric Variables')
print(numericVars.describe())

print('Summaries of Ordinal Variables')
print(ordinalVars.describe())
```

Summaries of Numeric Variables

	Lat	Lng	Population	Children	Age \
count	10000.000000	10000.000000	10000.000000	10000.0000	10000.000000
mean	38.757567	-90.782536	9756.562400	2.0877	53.078400
std	5.437389	15.156142	14432.698671	2.1472	20.698882
min	17.966120	-171.688150	0.000000	0.0000	18.000000
25%	35.341828	-97.082812	738.000000	0.0000	35.000000
50%	39.395800	-87.918800	2910.500000	1.0000	53.000000
75%	42.106908	-80.088745	13168.000000	3.0000	71.000000
max	70.640660	-65.667850	111850.000000	10.0000	89.000000

	Income	Outage_sec_perweek	Email	Contacts \
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	39806.926771	10.001848	12.016000	0.994200
std	28199.916702	2.976019	3.025898	0.988466
min	348.670000	0.099747	1.000000	0.000000
25%	19224.717500	8.018214	10.000000	0.000000
50%	33170.605000	10.018560	12.000000	1.000000
75%	53246.170000	11.969485	14.000000	2.000000
max	258900.700000	21.207230	23.000000	7.000000

	Yearly_equip_failure	Tenure	MonthlyCharge	Bandwidth_GB_Year
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	0.398000	34.526188	172.624816	3392.341550
std	0.635953	26.443063	42.943094	2185.294852
min	0.000000	1.000259	79.978860	155.506715
25%	0.000000	7.917694	139.979239	1236.470827
50%	0.000000	35.430507	167.484700	3279.536903
75%	1.000000	61.479795	200.734725	5586.141370
max	6.000000	71.999280	290.160419	7158.981530

Summaries of Ordinal Variables

Item1	Item2	Item3	Item4	Item5 \
-------	-------	-------	-------	---------

count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	3.490800	3.505100	3.487000	3.497500	3.492900
std	1.037797	1.034641	1.027977	1.025816	1.024819
min	1.000000	1.000000	1.000000	1.000000	1.000000
25%	3.000000	3.000000	3.000000	3.000000	3.000000
50%	3.000000	4.000000	3.000000	3.000000	3.000000
75%	4.000000	4.000000	4.000000	4.000000	4.000000
max	7.000000	7.000000	8.000000	7.000000	7.000000

	Item6	Item7	Item8
count	10000.000000	10000.000000	10000.000000
mean	3.497300	3.509500	3.495600
std	1.033586	1.028502	1.028633
min	1.000000	1.000000	1.000000
25%	3.000000	3.000000	3.000000
50%	3.000000	4.000000	3.000000
75%	4.000000	4.000000	4.000000
max	8.000000	7.000000	8.000000

[10]: *## C2 Summaries of Numeric Variables*

```
def quantDesc(data_frame, col_name):
    # Provide written description of the statistical summary output of .describe()

    count, mean, std, minimum, quarter, half, seventyfive, maximum =
    data_frame[col_name].describe()

    print(f'For the variable `{col_name}`:')
    print(f'There are {count} observations.')
    print(f'On average, the data tends towards the mean, which is {round(mean,
    2)}.')
    print(f'The standard deviation {round(std, 2)} is the amount of variation,
    or how much the data differs from the mean.')
    print(f'The smallest observation is {round(minimum, 2)}.')
    print(f'25% of the data falls below {round(quarter, 2)}.')
    print(f'50% of the data falls below {round(half, 2)}.')
    print(f'75% of the data falls below {round(seventyfive, 2)}.')
    print(f'100% of the data falls below the largest observation,
    {round(maximum, 2)}.')
    print('\n')
```

[11]: *## C2 Summaries of 13 Numeric Variables*

```
for col in numericVars:
    quantDesc(df_prep, col)
```

For the variable `Lat`:
There are 10000.0 observations.

On average, the data tends towards the mean, which is 38.76.
The standard deviation 5.44 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 17.97.
25% of the data falls below 35.34.
50% of the data falls below 39.4.
75% of the data falls below 42.11.
100% of the data falls below the largest observation, 70.64.

For the variable `Lng`:
There are 10000.0 observations.
On average, the data tends towards the mean, which is -90.78.
The standard deviation 15.16 is the amount of variation, or how much the data differs from the mean.
The smallest observation is -171.69.
25% of the data falls below -97.08.
50% of the data falls below -87.92.
75% of the data falls below -80.09.
100% of the data falls below the largest observation, -65.67.

For the variable `Population`:
There are 10000.0 observations.
On average, the data tends towards the mean, which is 9756.56.
The standard deviation 14432.7 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 0.0.
25% of the data falls below 738.0.
50% of the data falls below 2910.5.
75% of the data falls below 13168.0.
100% of the data falls below the largest observation, 111850.0.

For the variable `Children`:
There are 10000.0 observations.
On average, the data tends towards the mean, which is 2.09.
The standard deviation 2.15 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 0.0.
25% of the data falls below 0.0.
50% of the data falls below 1.0.
75% of the data falls below 3.0.
100% of the data falls below the largest observation, 10.0.

For the variable `Age`:
There are 10000.0 observations.

On average, the data tends towards the mean, which is 53.08.
The standard deviation 20.7 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 18.0.
25% of the data falls below 35.0.
50% of the data falls below 53.0.
75% of the data falls below 71.0.
100% of the data falls below the largest observation, 89.0.

For the variable `Income`:
There are 10000.0 observations.
On average, the data tends towards the mean, which is 39806.93.
The standard deviation 28199.92 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 348.67.
25% of the data falls below 19224.72.
50% of the data falls below 33170.6.
75% of the data falls below 53246.17.
100% of the data falls below the largest observation, 258900.7.

For the variable `Outage_sec_perweek`:
There are 10000.0 observations.
On average, the data tends towards the mean, which is 10.0.
The standard deviation 2.98 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 0.1.
25% of the data falls below 8.02.
50% of the data falls below 10.02.
75% of the data falls below 11.97.
100% of the data falls below the largest observation, 21.21.

For the variable `Email`:
There are 10000.0 observations.
On average, the data tends towards the mean, which is 12.02.
The standard deviation 3.03 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 1.0.
25% of the data falls below 10.0.
50% of the data falls below 12.0.
75% of the data falls below 14.0.
100% of the data falls below the largest observation, 23.0.

For the variable `Contacts`:
There are 10000.0 observations.

On average, the data tends towards the mean, which is 0.99.
The standard deviation 0.99 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 0.0.
25% of the data falls below 0.0.
50% of the data falls below 1.0.
75% of the data falls below 2.0.
100% of the data falls below the largest observation, 7.0.

For the variable `Yearly_equip_failure`:
There are 10000.0 observations.
On average, the data tends towards the mean, which is 0.4.
The standard deviation 0.64 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 0.0.
25% of the data falls below 0.0.
50% of the data falls below 0.0.
75% of the data falls below 1.0.
100% of the data falls below the largest observation, 6.0.

For the variable `Tenure`:
There are 10000.0 observations.
On average, the data tends towards the mean, which is 34.53.
The standard deviation 26.44 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 1.0.
25% of the data falls below 7.92.
50% of the data falls below 35.43.
75% of the data falls below 61.48.
100% of the data falls below the largest observation, 72.0.

For the variable `MonthlyCharge`:
There are 10000.0 observations.
On average, the data tends towards the mean, which is 172.62.
The standard deviation 42.94 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 79.98.
25% of the data falls below 139.98.
50% of the data falls below 167.48.
75% of the data falls below 200.73.
100% of the data falls below the largest observation, 290.16.

For the variable `Bandwidth_GB_Year`:
There are 10000.0 observations.

On average, the data tends towards the mean, which is 3392.34.
The standard deviation 2185.29 is the amount of variation, or how much the data differs from the mean.
The smallest observation is 155.51.
25% of the data falls below 1236.47.
50% of the data falls below 3279.54.
75% of the data falls below 5586.14.
100% of the data falls below the largest observation, 7158.98.

```
[12]: ## C2 Summaries of Binary Variables

def binarySummary(data_frame, col_name):
    # Get the counts, convert counts to percentages,
    # and calculate and display summary statistics

    counts = data_frame[col_name].value_counts()
    percentages = counts / counts.sum() * 100

    summary_stats = pd.DataFrame({'Count': counts, 'Percentage': percentages})
    summary_stats = summary_stats.rename(index={0: 'no', 1: 'yes'})

    print(f'Summary of `{col_name}`')
    print(summary_stats)
    print('\n')
```

```
[13]: ## C2 Summaries of 13 Binary Variables including Dependent Variable

print('Summary of dependent variable `Churn`')
binarySummary(df_prep, 'Churn')

for col in binaryVars:
    binarySummary(df, col)
```

Summary of dependent variable `Churn`

Summary of `Churn`

	Count	Percentage
Churn		
no	7350	73.5
yes	2650	26.5

Summary of `Techie`

	Count	Percentage
Techie		
no	8321	83.21
yes	1679	16.79

Summary of `Port_modem`

	Count	Percentage
Port_modem		
no	5166	51.66
yes	4834	48.34

Summary of `Tablet`

	Count	Percentage
Tablet		
no	7009	70.09
yes	2991	29.91

Summary of `Phone`

	Count	Percentage
Phone		
yes	9067	90.67
no	933	9.33

Summary of `Multiple`

	Count	Percentage
Multiple		
no	5392	53.92
yes	4608	46.08

Summary of `OnlineSecurity`

	Count	Percentage
OnlineSecurity		
no	6424	64.24
yes	3576	35.76

Summary of `OnlineBackup`

	Count	Percentage
OnlineBackup		
no	5494	54.94
yes	4506	45.06

Summary of `DeviceProtection`

	Count	Percentage
DeviceProtection		
no	5614	56.14

yes	4386	43.86
-----	------	-------

Summary of `TechSupport`

	Count	Percentage
TechSupport		
no	6250	62.5
yes	3750	37.5

Summary of `StreamingTV`

	Count	Percentage
StreamingTV		
no	5071	50.71
yes	4929	49.29

Summary of `StreamingMovies`

	Count	Percentage
StreamingMovies		
no	5110	51.1
yes	4890	48.9

Summary of `PaperlessBilling`

	Count	Percentage
PaperlessBilling		
yes	5882	58.82
no	4118	41.18

```
[14]: ## C2 Summaries of Categorical Variables
```

```
def catSummary(data_frame, col_name):  
    # Get the counts, convert counts to percentages,  
    # and calculate and display summary statistics  
  
    counts = data_frame[col_name].value_counts()  
    percentages = counts / counts.sum() * 100  
  
    summary_stats = pd.DataFrame({'Count': counts, 'Percentage': percentages})  
  
    print(f'Summary of `{col_name}`')  
    print(summary_stats)  
    print('\n')
```

```
[15]: ## C2 Summaries

# Summaries of Six Categorical Variables
for col in catVars:
    catSummary(df, col)

# Summaries of Eight Ordinal Variables
for col in ordinalVars:
    catSummary(df_prep, col)
```

Summary of `Area`

	Count	Percentage
Area		
Suburban	3346	33.46
Urban	3327	33.27
Rural	3327	33.27

Summary of `Marital`

	Count	Percentage
Marital		
Divorced	2092	20.92
Widowed	2027	20.27
Separated	2014	20.14
Never Married	1956	19.56
Married	1911	19.11

Summary of `Gender`

	Count	Percentage
Gender		
Female	5025	50.25
Male	4744	47.44
Nonbinary	231	2.31

Summary of `Contract`

	Count	Percentage
Contract		
Month-to-month	5456	54.56
Two Year	2442	24.42
One year	2102	21.02

Summary of `InternetService`

	Count	Percentage
InternetService		
Fiber Optic	4408	44.08

DSL	3463	34.63
None	2129	21.29

Summary of `PaymentMethod`

	Count	Percentage
PaymentMethod		
Electronic Check	3398	33.98
Mailed Check	2290	22.90
Bank Transfer(automatic)	2229	22.29
Credit Card (automatic)	2083	20.83

Summary of `Item1`

	Count	Percentage
Item1		
3	3448	34.48
4	3358	33.58
2	1393	13.93
5	1359	13.59
1	224	2.24
6	199	1.99
7	19	0.19

Summary of `Item2`

	Count	Percentage
Item2		
3	3415	34.15
4	3412	34.12
5	1368	13.68
2	1360	13.60
1	217	2.17
6	215	2.15
7	13	0.13

Summary of `Item3`

	Count	Percentage
Item3		
3	3435	34.35
4	3410	34.10
2	1424	14.24
5	1313	13.13
6	203	2.03
1	202	2.02
7	12	0.12
8	1	0.01

Summary of `Item4`

	Count	Percentage
Item4		
4	3452	34.52
3	3430	34.30
2	1350	13.50
5	1335	13.35
1	221	2.21
6	203	2.03
7	9	0.09

Summary of `Item5`

	Count	Percentage
Item5		
3	3462	34.62
4	3417	34.17
2	1378	13.78
5	1321	13.21
1	206	2.06
6	204	2.04
7	12	0.12

Summary of `Item6`

	Count	Percentage
Item6		
3	3445	34.45
4	3333	33.33
2	1427	14.27
5	1382	13.82
6	210	2.10
1	190	1.90
7	12	0.12
8	1	0.01

Summary of `Item7`

	Count	Percentage
Item7		
4	3456	34.56
3	3446	34.46
5	1335	13.35
2	1309	13.09
6	224	2.24
1	219	2.19


```
7          11          0.11
```

Summary of `Item8`

	Count	Percentage
Item8		
3	3461	34.61
4	3400	34.00
2	1378	13.78
5	1335	13.35
1	206	2.06
6	205	2.05
7	14	0.14
8	1	0.01

3.1.3 C3. Univariate and Bivariate Visualizations

The code below generates the univariate and bivariate visualizations of the distributions of the dependent and independent variables. The bivariate visualizations consider `Churn` as the dependent variable.

```
[16]: ## C3 Univariate Visualizations

import matplotlib.pyplot as plt

def plot_histogram(data_frame, col_name):
    # Overlays a transparent boxplot over a histogram
    data = data_frame[col_name].values
    fig, ax1 = plt.subplots()

    ax1.hist(data, bins = 10, alpha = 0.7, label = 'Histogram')
    ax2 = ax1.twinx()
    ax2.boxplot(data, vert = False, widths = 0.5, patch_artist = True,
                boxprops = dict(facecolor = 'orange', alpha = 0.5))

    ax1.set_ylabel('Frequency')
    ax2.set_ylabel('Boxplot')

    plt.title(f'Histogram with Boxplot for {col_name}')
    plt.show()

def plot_binaryhist(data_frame, col_name):
    # Histogram of a binary variable
    catCounts = data_frame[col_name].value_counts()
```

```

cats = catCounts.index.tolist()
counts = catCounts.values.tolist()

plt.bar(cats, counts)
plt.xticks(ticks = [0, 1], labels = ['No', 'Yes'])
plt.title(f'Histogram of {col_name}')
plt.xlabel('Category')
plt.ylabel('Frequency')
plt.show()

def plot_cathist(data_frame, col_name):
    # Histogram of a categorical variable
    catCounts = data_frame[col_name].value_counts()
    cats = catCounts.index.tolist()
    counts = catCounts.values.tolist()

    plt.bar(cats, counts)
    plt.title(f'Histogram of {col_name}')
    plt.xlabel('Category')
    plt.ylabel('Frequency')
    plt.show()

```

```

[17]: ## C3 Univariate Visualizations

# Histograms + Boxplots for the dependent variable
plot_binaryhist(df_prep, 'Churn')

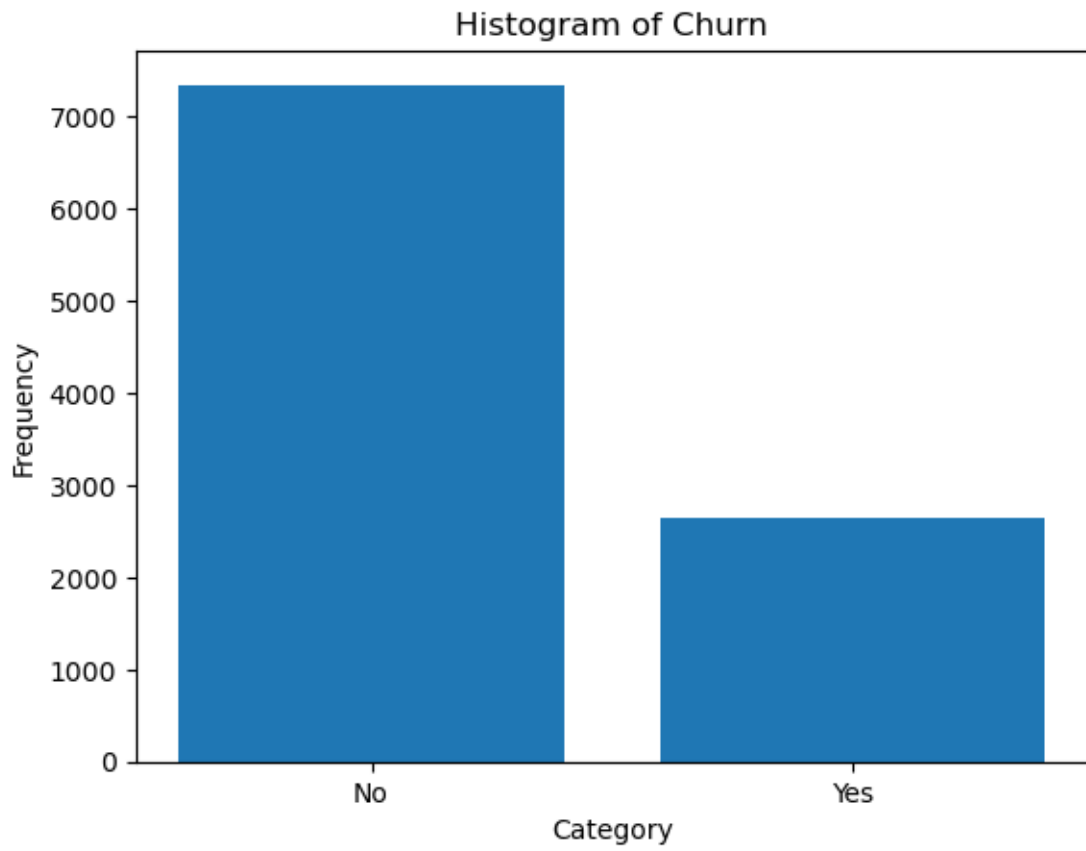
# Histograms + Boxplots for 13 Numeric Variables
for col in numericVars:
    plot_histobox(df_prep, col)

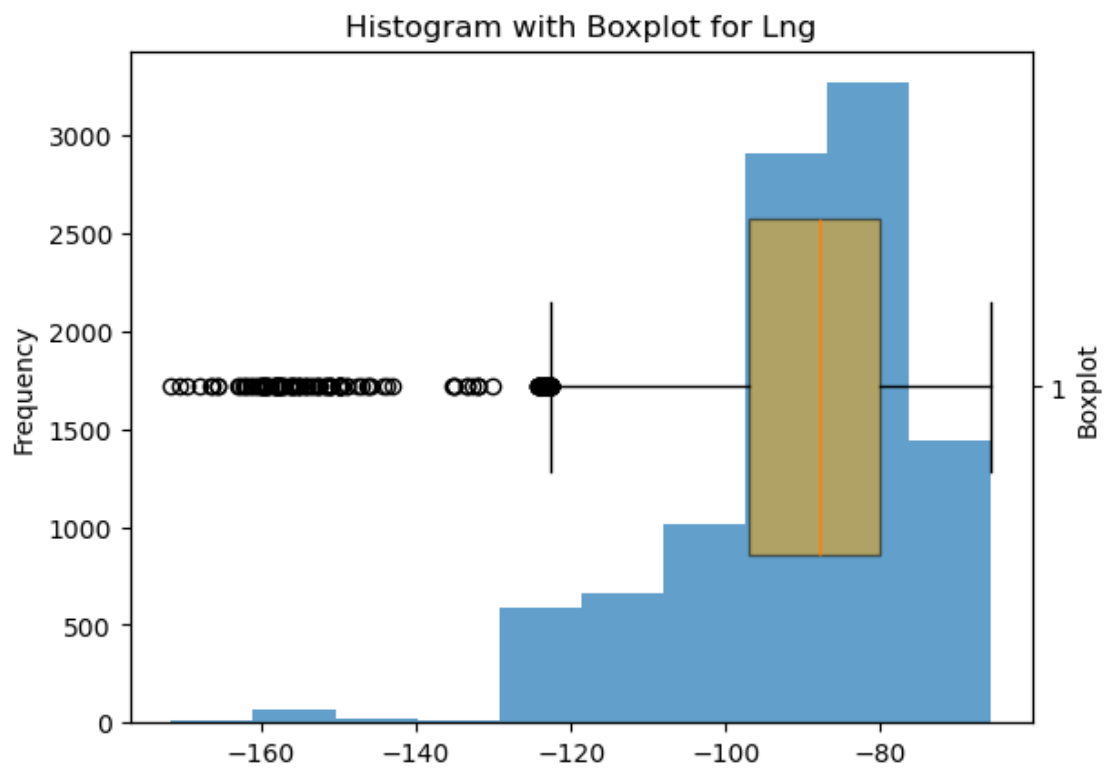
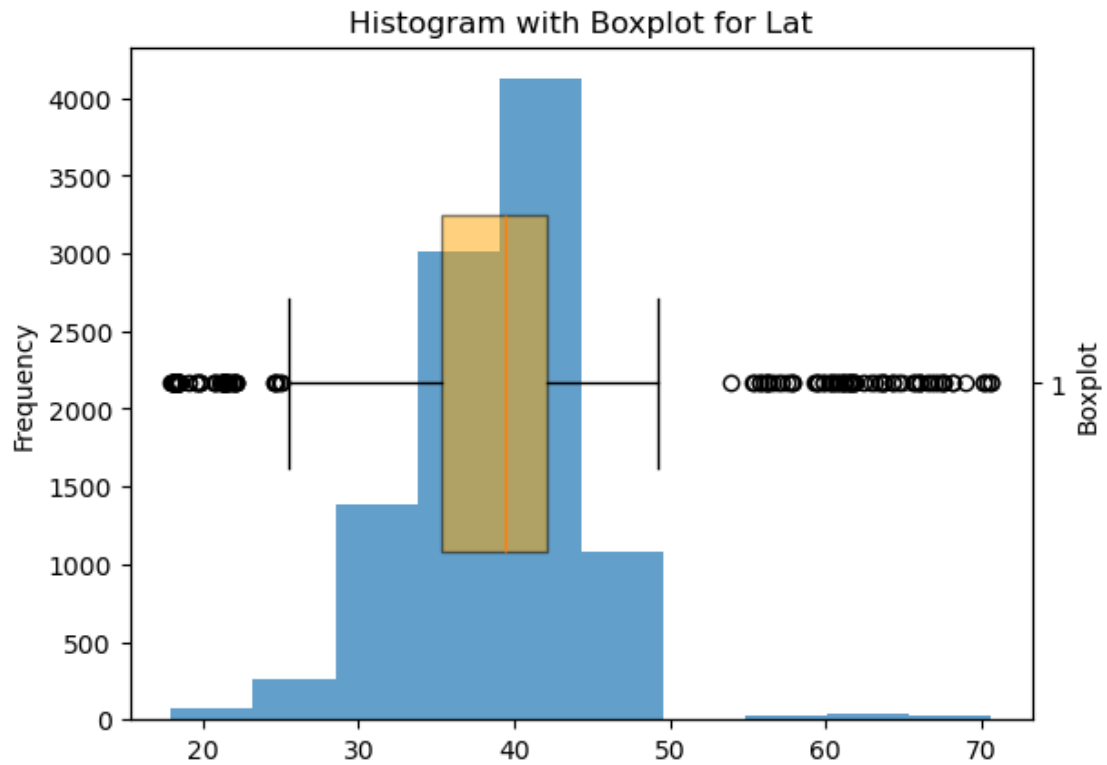
# Histograms for 12 Binary Variables
for col in binaryVars:
    plot_binaryhist(df_prep, col)


# Histograms for 6 Categorical Variables
for col in catVars:
    plot_cathist(df, col)


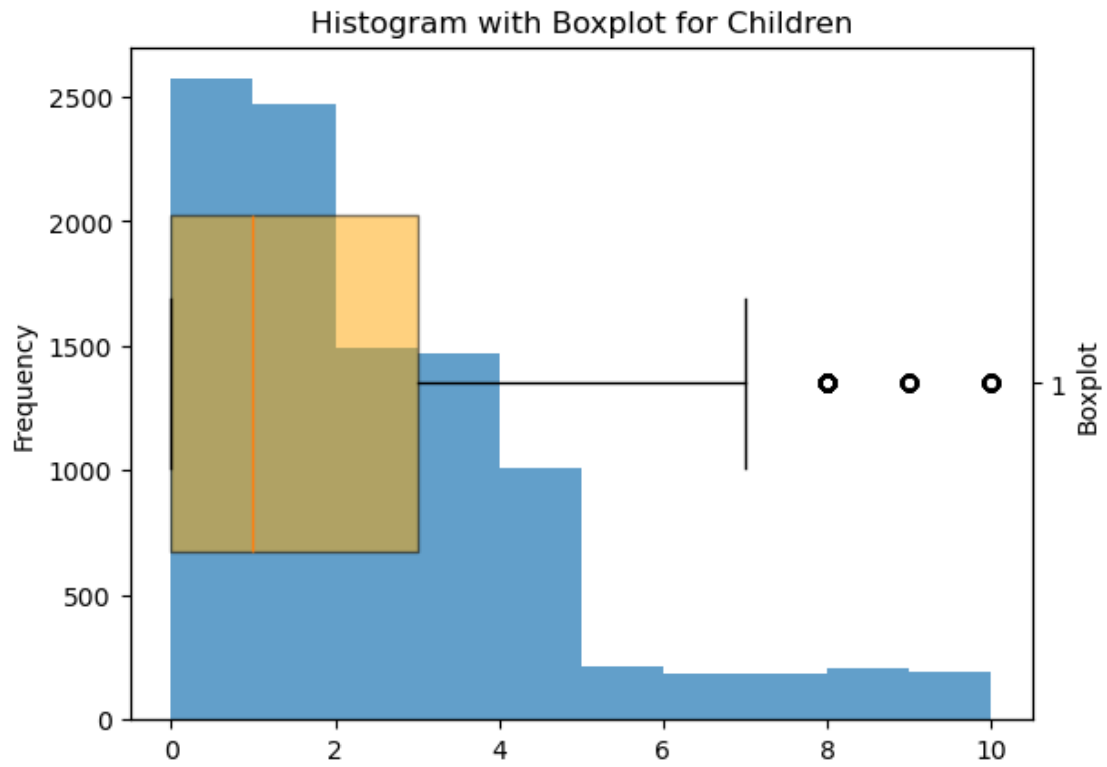
# Histograms for 8 Ordinal Variables
for col in ordinalVars:
    plot_cathist(df_prep, col)

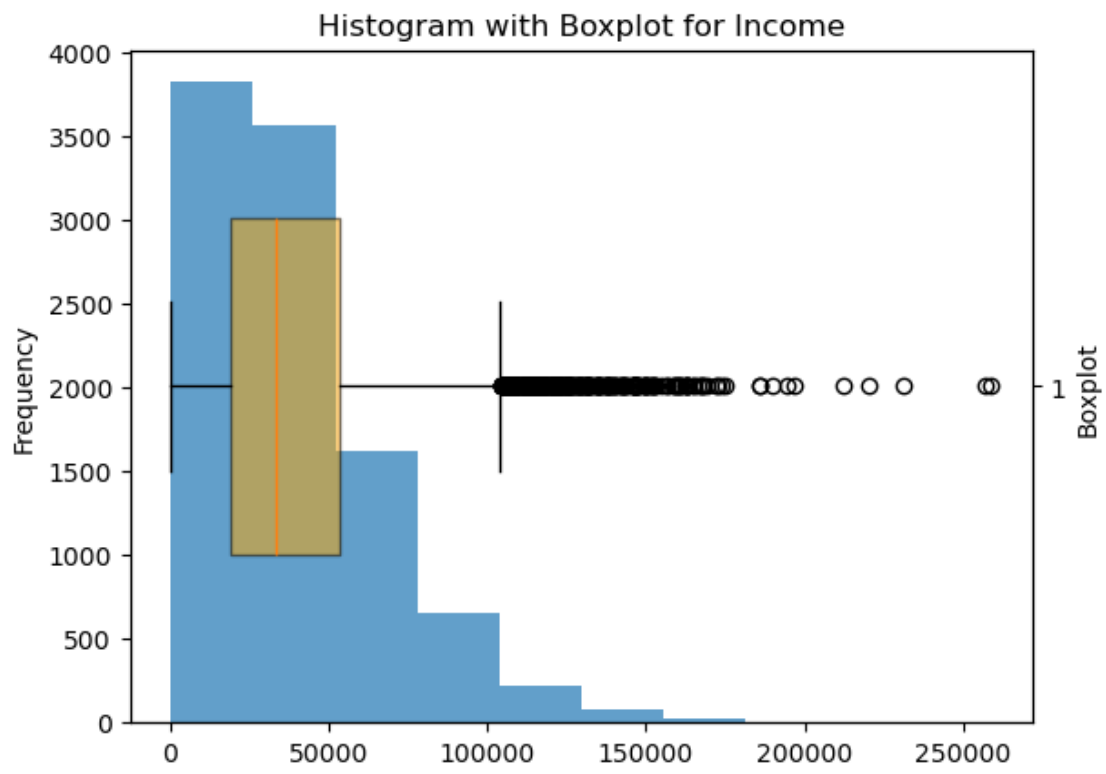
```

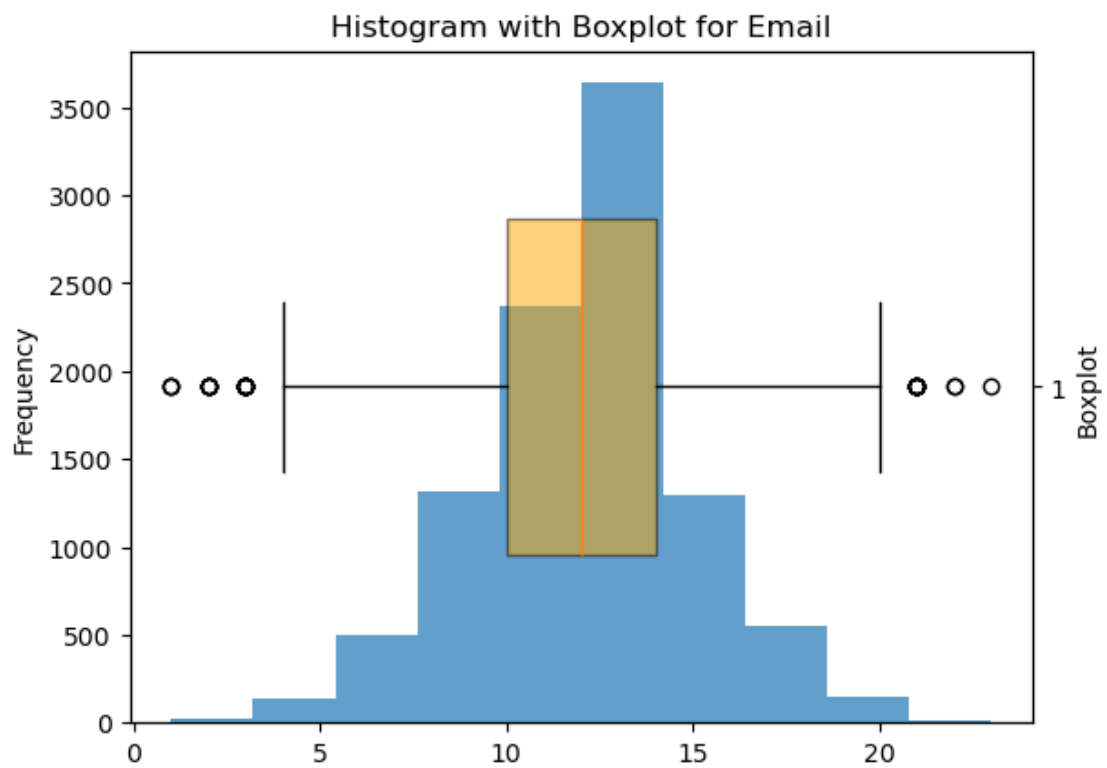



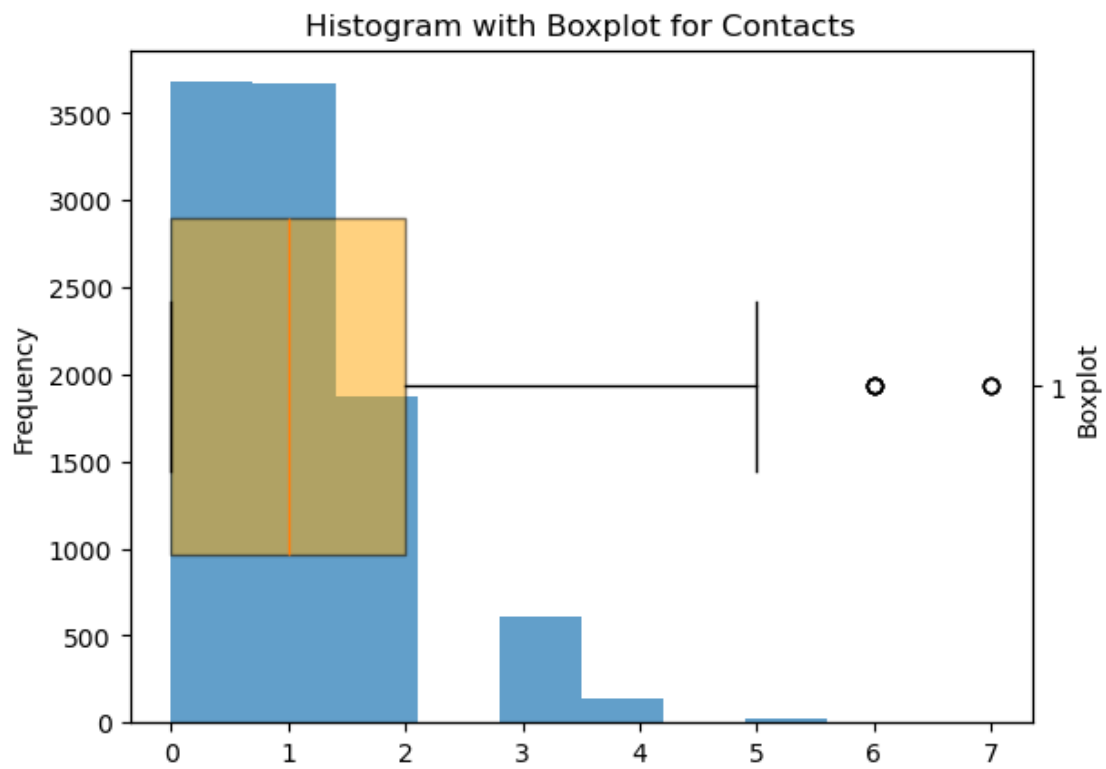


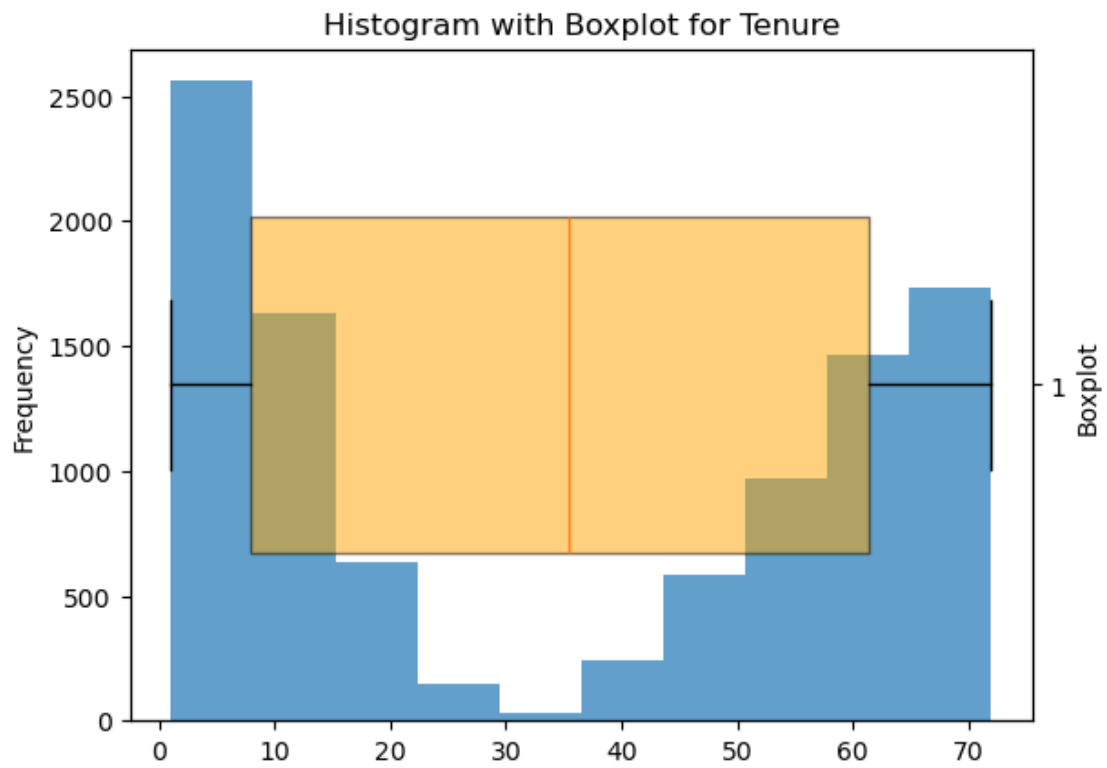
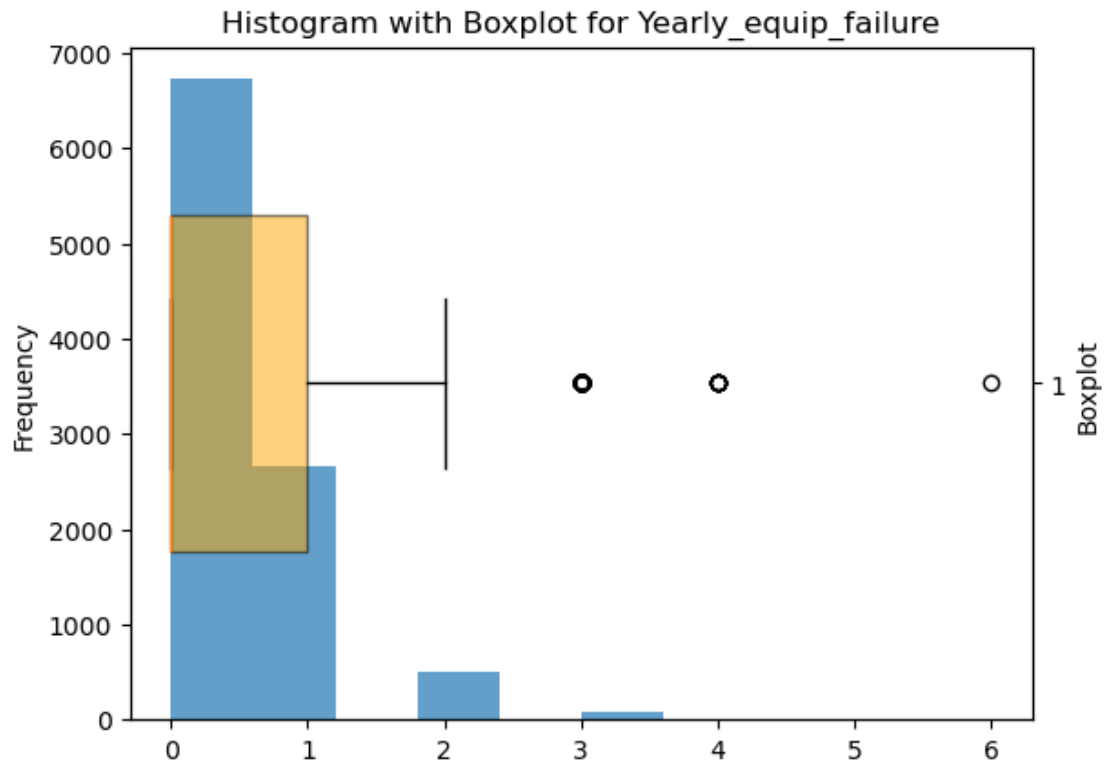








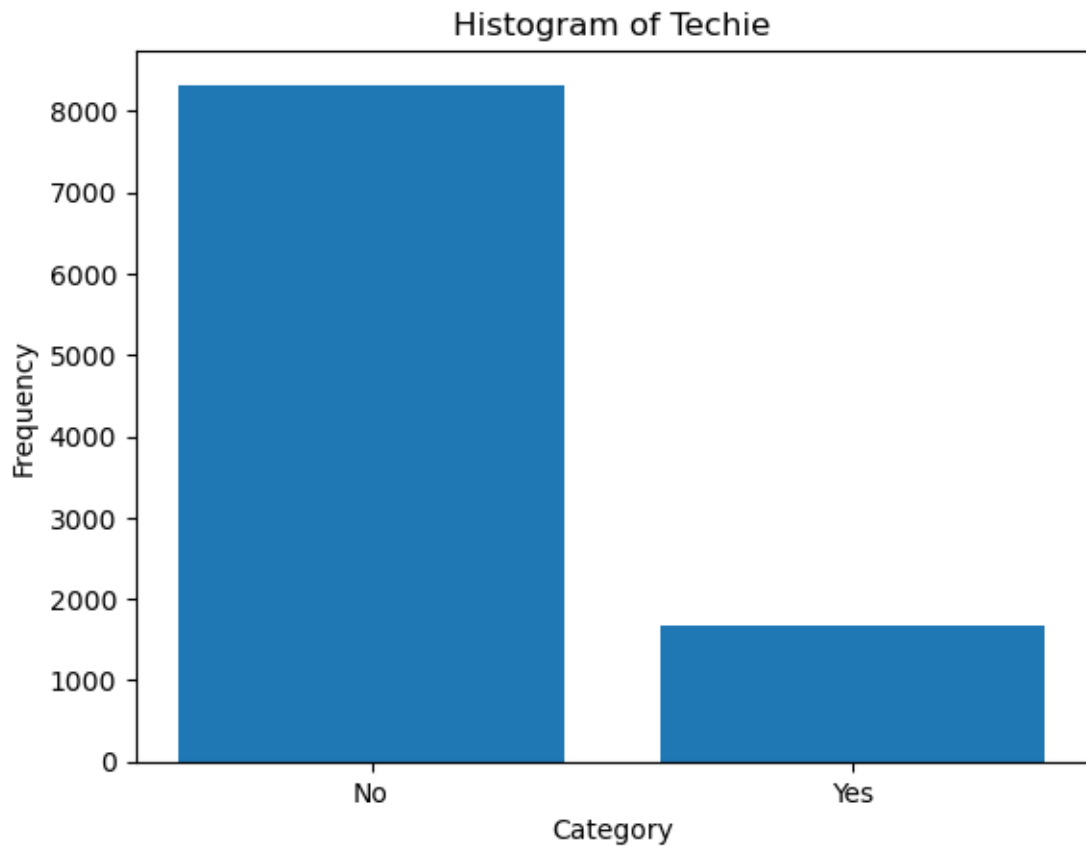





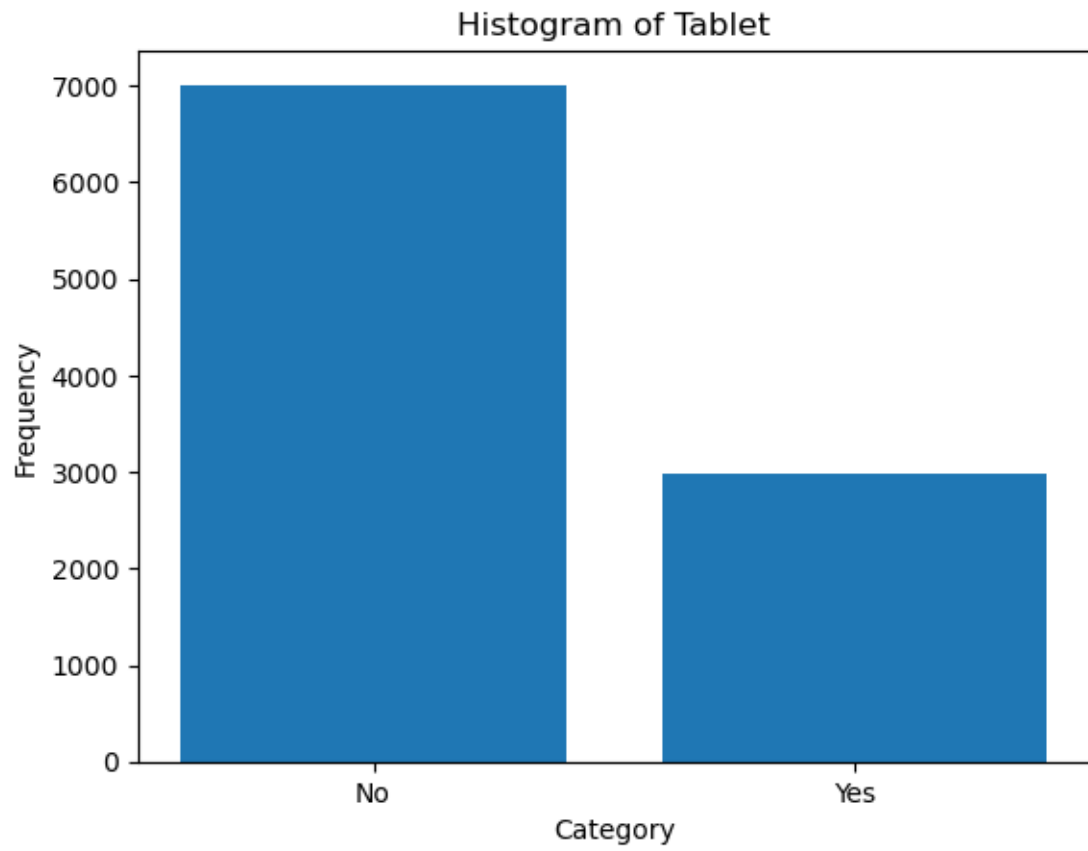





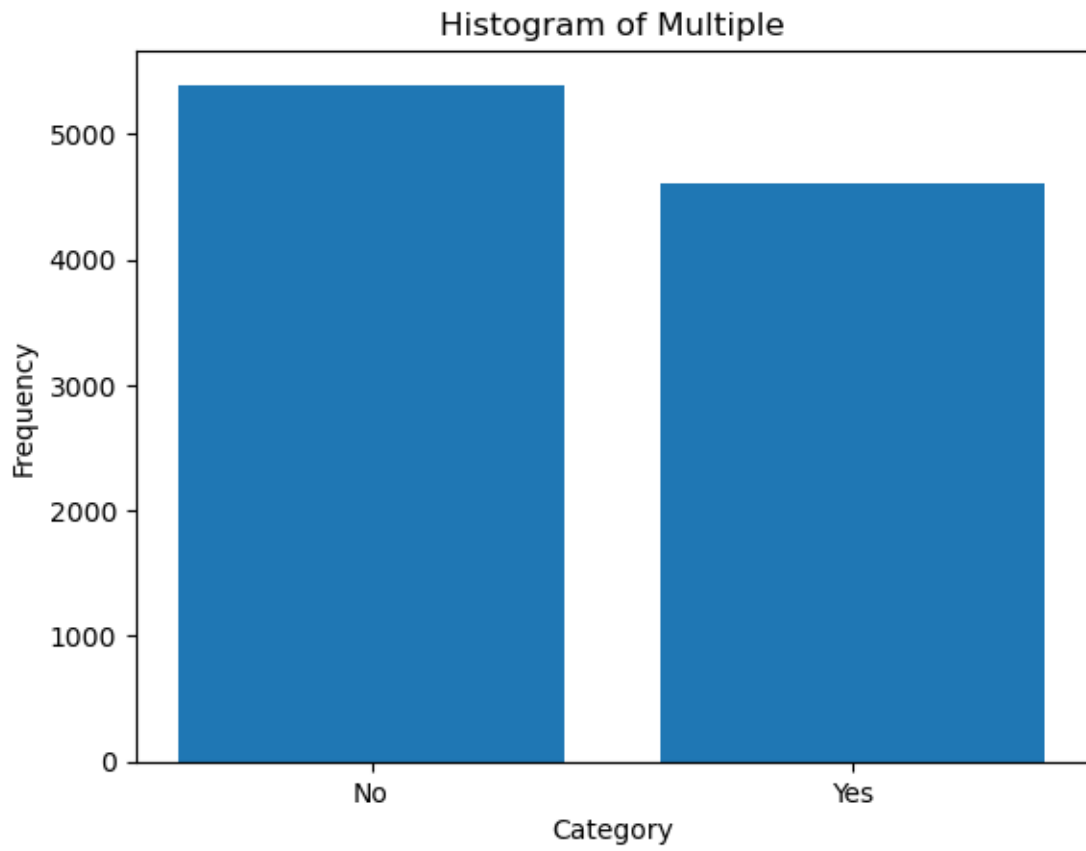


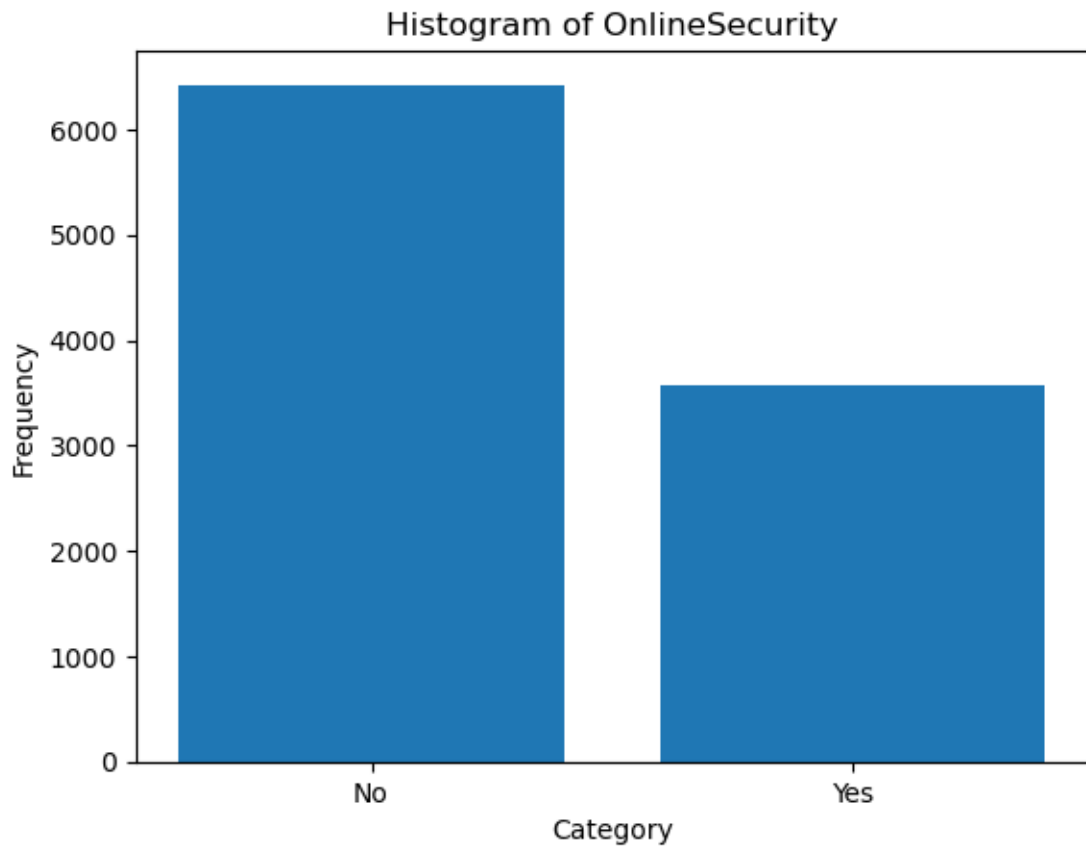








































































```
[18]: ## C3 Bivariate Visualizations

import seaborn as sns
import warnings
warnings.filterwarnings("ignore", category = FutureWarning, module = "seaborn")

def plot_bivarcath(data_frame, cont_col, cat_col):
    # Plot bivariate boxplots of categorical vs continuous variables
    sns.boxplot(x = df['Income'], y = df['Churn'])
    plt.title(f'Box Plot of {cat_col} by {cont_col}')
    plt.xlabel(cont_col)
    plt.ylabel(cat_col)
    plt.show()

def plot_bivarBars(data_frame, exp, dep):
    # Plot bivariate stacked bar plots of two categorical variables
    counts = pd.crosstab(data_frame[exp], data_frame[dep])
    counts.plot(kind = 'bar', stacked = True)

    plt.title(f'Stacked Bar Charts of `{dep}` by `{exp}`')
```

```

plt.xticks(rotation = 0)
plt.xlabel(exp)
plt.ylabel(dep)
plt.show()

def plot_stackedBars(data_frame, exp, dep):
    # Stacked bar charts of categorical vs ordinal
    sns.countplot(data = data_frame, x = exp, hue = dep)
    plt.title(f'Grouped Bar Charts of `{dep}` by `{exp}`')
    plt.xticks(rotation = 0)
    plt.xlabel(exp)
    plt.ylabel(dep)
    plt.show()

```

```

[19]: ## C3 Bivariate Visualizations


# Dependent Variable vs 13 Numeric Variables
for col in numericVars:
    plot_bivarcath(df, col, 'Churn')


# Dependent Variable vs 12 Categorical Variables
for col in binaryVars:
    plot_bivarBars(df, col, 'Churn')


# Dependent Variable vs 6 Categorical Variables
for col in catVars:
    plot_bivarBars(df, col, 'Churn')


# Dependent Variable vs 8 Ordinal Variables
for col in ordinalVars:
    plot_stackedBars(df, col, 'ChurnStr')


```

















































































































3.1.4 C4. Description of Data Transformation Goals and Steps to Achieve Goals

The data wrangling performed on the data set consisted of the following:

- Re-expression of the thirteen binary variables by encoding True and False as 1 and 0, respectively. This was done using a for loop that applies `.replace()` to specific columns using a dictionary. It was also necessary to type cast these with `.astype('category')` in order for the bivariate graphs to generate. The thirteen re-expressed binary variables:
 - Churn, Techie, Port_modem, Tablet, Phone, Multiple, OnlineSecurity, OnlineBackup, DeviceProtection, TechSupport, StreamingTV, StreamingMovies, PaperlessBilling
- Six categorical variable were re-expressed using one-hot encoding with `.get_dummies()` from Pandas. It was also necessary to type cast these with `.astype('category')` in order for the bivariate graphs to generate. The six re-expressed categorical variables:
 - Area, Marital, Gender, Contract, InternetService, PaymentMethod

See code attached, in D208_PA_MendezD_Task2_Attempt2.ipynb.

3.1.5 C5. Prepared Data Set as CSV file

```
[20]: ## C5 CSV Output

# df_prep.to_csv('D208_PA_MendezD_Task2_Attempt1.csv', sep = ',', encoding = 'utf-8', index = False)
```

4 Part IV: Model Comparison and Analysis

4.1 D. Comparison of Initial and Reduced Logistic Regression Models

4.1.1 D1. Initial Logistic Regression Model

The following cells contain the code necessary to construct the initial logistic regression model with all 39 variables identified in C2. The summary of the initial model is printed below.

```
[21]: ## D1 Initial Logistic Regression Model with 39 independent variables

import statsmodels.api as sm

# Assign 'Churn' as the dependent variable
depVar = df_prep['Churn']

# Create explanatory variable data frame
expVars = df_prep.drop('Churn', axis = 1)

# Add intercept to the model
expVars_intercept = sm.add_constant(expVars)

# Fit the intercept model
model_intercept = sm.Logit(depVar, expVars_intercept).fit()

# Print the model summary
print(model_intercept.summary())
```

Optimization terminated successfully.

Current function value: 0.217086

Iterations 9

Logit Regression Results

```
=====
Dep. Variable:          Churn    No. Observations:          10000
Model:                Logit    Df Residuals:              9951
Method:                MLE     Df Model:                  48
Date:                  Sun, 07 Jul 2024    Pseudo R-squ.:          0.6246
Time:                  12:04:48    Log-Likelihood:         -2170.9
converged:              True    LL-Null:              -5782.2
Covariance Type:        nonrobust    LLR p-value:            0.000
=====
```

			coef	std err	z
P> z	[0.025	0.975]			

const			-5.3777	1.546	-3.479
0.001	-8.408	-2.348			
Lat			0.0032	0.007	0.452
0.651	-0.011	0.017			
Lng			-0.0019	0.003	-0.735
0.462	-0.007	0.003			
Population			-1.64e-07	2.8e-06	-0.059
0.953	-5.65e-06	5.32e-06			
Children			-0.0055	0.137	-0.040
0.968	-0.274	0.263			
Age			0.0028	0.015	0.194
0.846	-0.026	0.032			
Income			3.994e-07	1.38e-06	0.290
0.772	-2.3e-06	3.1e-06			
Outage_sec_perweek			-0.0027	0.013	-0.209
0.835	-0.028	0.023			
Email			-0.0094	0.013	-0.738
0.460	-0.034	0.016			
Contacts			0.0626	0.039	1.606
0.108	-0.014	0.139			
Yearly equip_failure			-0.0348	0.061	-0.569
0.569	-0.155	0.085			
Tenure			-0.1685	0.363	-0.464
0.642	-0.879	0.542			
MonthlyCharge			0.0394	0.014	2.851
0.004	0.012	0.067			
Bandwidth_GB_Year			0.0006	0.004	0.143
0.886	-0.008	0.009			
Techie			1.0981	0.103	10.672
0.000	0.896	1.300			
Port_modem			0.1435	0.077	1.853
0.064	-0.008	0.295			
Tablet			-0.0508	0.085	-0.600
0.549	-0.217	0.115			
Phone			-0.2988	0.133	-2.246
0.025	-0.560	-0.038			
Multiple			0.3602	0.202	1.786
0.074	-0.035	0.755			
OnlineSecurity			-0.2935	0.312	-0.942
0.346	-0.905	0.317			
OnlineBackup			-0.1172	0.181	-0.649
0.517	-0.471	0.237			
DeviceProtection			-0.1085	0.234	-0.464
0.642	-0.566	0.349			

TechSupport			-0.2144	0.173	-1.237
0.216	-0.554	0.125			
StreamingTV			1.1140	0.510	2.185
0.029	0.115	2.113			
StreamingMovies			1.2744	0.363	3.506
0.000	0.562	1.987			
PaperlessBilling			0.1661	0.079	2.108
0.035	0.012	0.320			
Area_Suburban			-0.0481	0.096	-0.503
0.615	-0.235	0.139			
Area_Urban			0.0517	0.095	0.547
0.584	-0.134	0.237			
Marital_Married			0.1045	0.122	0.856
0.392	-0.135	0.344			
Marital_Never Married			0.0133	0.122	0.109
0.913	-0.226	0.253			
Marital_Separated			0.1179	0.120	0.979
0.327	-0.118	0.354			
Marital_Widowed			0.2603	0.121	2.158
0.031	0.024	0.497			
Gender_Male			0.2276	0.288	0.790
0.430	-0.337	0.793			
Gender_Nonbinary			-0.0775	0.282	-0.275
0.784	-0.631	0.476			
Contract_One year			-3.4144	0.129	-26.542
0.000	-3.667	-3.162			
Contract_Two Year			-3.5185	0.127	-27.725
0.000	-3.767	-3.270			
InternetService_Fiber Optic			-1.9156	2.095	-0.914
0.361	-6.022	2.191			
InternetService_None			-0.7259	1.671	-0.434
0.664	-4.000	2.549			
PaymentMethod_Credit Card (automatic)			0.2094	0.118	1.774
0.076	-0.022	0.441			
PaymentMethod_Electronic Check			0.6308	0.106	5.946
0.000	0.423	0.839			
PaymentMethod_Mailed Check			0.2392	0.116	2.056
0.040	0.011	0.467			
Item1			-0.0189	0.055	-0.346
0.730	-0.126	0.088			
Item2			-0.0040	0.052	-0.077
0.939	-0.106	0.098			
Item3			0.0206	0.047	0.437
0.662	-0.072	0.113			
Item4			-0.0302	0.042	-0.718
0.473	-0.113	0.052			
Item5			-0.0362	0.044	-0.815
0.415	-0.123	0.051			

Item6			-0.0193	0.045	-0.427
0.670	-0.108	0.069			
Item7			-0.0026	0.043	-0.061
0.951	-0.087	0.082			
Item8			-0.0141	0.040	-0.348
0.728	-0.093	0.065			

=====

=====

4.1.2 D2. Model Reduction Method and Justification

The model reduction method utilized below is backward stepwise elimination. Beginning with all explanatory variables, backwards elimination removes the least significant variable, as determined by p-value, at each iteration of the loop. This is iterated until no improvement is observed, or rather, no variables are observed to have $p > 0.05$ (Middleton, 2022).

```
[22]: # D2 Model Reduction Method

def backward_elimination_Logit(data, target, siglevel):
    # Function sourced from AnalyticsVidhya to perform Backwards Elimination
    # adapted for Logistic Regression

    features = data.columns.tolist()
    while(len(features) > 1):
        features_with_constant = sm.add_constant(data[features])
        model = sm.Logit(target, features_with_constant)
        result = model.fit()

        p_values = result.pvalues[1:]
        max_p_value = p_values.max()

        if (max_p_value >= siglevel):
            excluded_feature = p_values.idxmax()
            features.remove(excluded_feature)
        else:
            break
    return features, result
```

```
[23]: # D2 Model Reduction Method

selected_features, model_reduced = backward_elimination_Logit(expVars, depVar,
    ↪0.05)

print(selected_features)
```

Optimization terminated successfully.
Current function value: 0.217086

Iterations 9
Optimization terminated successfully.
Current function value: 0.217086
Iterations 9
Optimization terminated successfully.
Current function value: 0.217086
Iterations 9
Optimization terminated successfully.
Current function value: 0.217086
Iterations 9
Optimization terminated successfully.
Current function value: 0.217086
Iterations 9
Optimization terminated successfully.
Current function value: 0.217087
Iterations 9
Optimization terminated successfully.
Current function value: 0.217089
Iterations 9
Optimization terminated successfully.
Current function value: 0.217093
Iterations 9
Optimization terminated successfully.
Current function value: 0.217098
Iterations 9
Optimization terminated successfully.
Current function value: 0.217105
Iterations 9
Optimization terminated successfully.
Current function value: 0.217114
Iterations 9
Optimization terminated successfully.
Current function value: 0.217120
Iterations 9
Optimization terminated successfully.
Current function value: 0.217131
Iterations 9
Optimization terminated successfully.
Current function value: 0.217143
Iterations 9
Optimization terminated successfully.
Current function value: 0.217161
Iterations 9
Optimization terminated successfully.
Current function value: 0.217181
Iterations 9
Optimization terminated successfully.
Current function value: 0.217204

Iterations 9
Optimization terminated successfully.
Current function value: 0.217230
Iterations 9
Optimization terminated successfully.
Current function value: 0.217251
Iterations 9
Optimization terminated successfully.
Current function value: 0.217269
Iterations 9
Optimization terminated successfully.
Current function value: 0.217299
Iterations 9
Optimization terminated successfully.
Current function value: 0.217306
Iterations 9
Optimization terminated successfully.
Current function value: 0.217336
Iterations 9
Optimization terminated successfully.
Current function value: 0.217347
Iterations 9
Optimization terminated successfully.
Current function value: 0.217384
Iterations 9
Optimization terminated successfully.
Current function value: 0.217424
Iterations 9
Optimization terminated successfully.
Current function value: 0.217465
Iterations 9
Optimization terminated successfully.
Current function value: 0.217510
Iterations 9
Optimization terminated successfully.
Current function value: 0.217631
Iterations 9
Optimization terminated successfully.
Current function value: 0.217787
Iterations 9
Optimization terminated successfully.
Current function value: 0.217875
Iterations 9
Optimization terminated successfully.
Current function value: 0.218031
Iterations 9
['Tenure', 'MonthlyCharge', 'Techie', 'Phone', 'Multiple', 'OnlineSecurity',
'TechSupport', 'StreamingTV', 'StreamingMovies', 'PaperlessBilling',

```
'Marital_Widowed', 'Gender_Male', 'Contract_One year', 'Contract_Two Year',
'InternetService_Fiber Optic', 'InternetService_None', 'PaymentMethod_Electronic
Check']
```

4.1.3 D3. Reduced Model

Upon applying Backwards Elimination, the explanatory variables of the reduced model are:

- Tenure
- MonthlyCharge
- Techie
- Phone
- Multiple
- OnlineSecurity
- TechSupport
- StreamingTV
- StreamingMovies
- PaperlessBilling
- Marital_Widowed
- Gender_Male
- Contract_One year
- Contract_Two Year
- InternetService_Fiber Optic
- InternetService_None
- PaymentMethod_Electronic Check

The corresponding reduced model was generated above in the function `backward_elimination_Logit()`. Below is the code that prints its summary.

```
[24]: # D3 Reduced Model

print(model_reduced.summary())
```

```

                                Logit Regression Results
=====
Dep. Variable:                  Churn    No. Observations:                  10000
Model:                          Logit    Df Residuals:                      9982
Method:                          MLE     Df Model:                          17
Date:                            Sun, 07 Jul 2024    Pseudo R-squ.:                  0.6229
Time:                            12:08:14    Log-Likelihood:                 -2180.3
converged:                       True     LL-Null:                         -5782.2
Covariance Type:                 nonrobust    LLR p-value:                    0.000
=====
=====
                                coef    std err          z      P>|z|
-----
[0.025    0.975]
-----
const                -4.7283      0.314    -15.042      0.000
```

-5.344	-4.112				
Tenure		-0.1158	0.003	-39.343	0.000
-0.122	-0.110				
MonthlyCharge		0.0375	0.003	14.301	0.000
0.032	0.043				
Techie		1.0987	0.102	10.725	0.000
0.898	1.300				
Phone		-0.2898	0.132	-2.197	0.028
-0.548	-0.031				
Multiple		0.4516	0.112	4.049	0.000
0.233	0.670				
OnlineSecurity		-0.2388	0.081	-2.967	0.003
-0.397	-0.081				
TechSupport		-0.1907	0.085	-2.251	0.024
-0.357	-0.025				
StreamingTV		1.3345	0.137	9.735	0.000
1.066	1.603				
StreamingMovies		1.5051	0.157	9.582	0.000
1.197	1.813				
PaperlessBilling		0.1633	0.078	2.085	0.037
0.010	0.317				
Marital_Widowed		0.1927	0.096	2.006	0.045
0.004	0.381				
Gender_Male		0.2759	0.077	3.570	0.000
0.124	0.427				
Contract_One year		-3.4023	0.128	-26.665	0.000
-3.652	-3.152				
Contract_Two Year		-3.4998	0.126	-27.808	0.000
-3.747	-3.253				
InternetService_Fiber Optic		-2.1360	0.110	-19.445	0.000
-2.351	-1.921				
InternetService_None		-0.9998	0.113	-8.825	0.000
-1.222	-0.778				
PaymentMethod_Electronic Check		0.4759	0.082	5.833	0.000
0.316	0.636				

=====

=====

4.2 E. Analysis Using Reduced Logistic Regression Model

4.2.1 E1. Model Comparison using Model Evaluation Metrics

The initial model can be compared with the reduced model using the Pseudo R^2 and AIC. The following comparisons can be drawn using the output of the code below.

- Pseudo R^2 : The slight decrease in Pseudo R^2 from the initial model to the reduced model suggests that the initial model has slightly better fit (FAQ: What are pseudo R-squareds?).
- AIC: Since the reduced model has a higher AIC, it has a worse fit relative to the initial model.

```
[25]: ## E1 Model Comparison

def modelCompare(initialModel, reducedModel):
    # Prints a data frame as a table that compares model evaluation metrics for two
    ↪ regression models

    evalMetrics = {
        'Criteria': ['Pseudo R2', 'AIC'],
        'Initial': [initialModel.prsquared, initialModel.aic],
        'Reduced': [reducedModel.prsquared, reducedModel.aic]
    }

    em = pd.DataFrame(evalMetrics)
    print(em)
```

```
[26]: ## E1 Model Comparison

modelCompare(model_intercept, model_reduced)
```

	Criteria	Initial	Reduced
0	Pseudo R2	0.624564	0.622929
1	AIC	4439.713102	4396.618743

```
[27]: ## E1 Model Comparison
# Sourced from Statology, Bobbitt, 2021

import scipy

def likelihoodRatioTest(full_model, reduced_model, sig_level):
    # Computes the likelihood ratio test statistic for two regression models
    lr_stat = -2 * (reduced_model.llf - full_model.llf)
    df_diff = full_model.df_model - reduced_model.df_model
    p_value = scipy.stats.chi2.sf(lr_stat, df_diff)

    print(f"Likelihood Ratio Test Statistic: {lr_stat}")
    if p_value < sig_level:
        print('Reject null hypothesis. Evidence to conclude that the full model
        ↪ offers a significantly better fit.')
    if p_value >= sig_level:
        print('Fail to reject null hypothesis. Evidence to conclude that the
        ↪ nested model offers a significantly better fit.')
```

```
[28]: ## E1 Model Comparison

likelihoodRatioTest(model_intercept, model_reduced, 0.05)
```

Likelihood Ratio Test Statistic: 18.905641219661447
 Fail to reject null hypothesis. Evidence to conclude that the nested model

offers a significantly better fit.

4.2.2 E2. Confusion Matrix and Accuracy Calculation

Below is the code that generates the confusion matrix and performs the accuracy calculation.

```
[29]: ## E2 Confusion Matrix and Accuracy Calculation
# Code sourced from Geeksforgeeks

# Data frame of explanatory test data
expTest = df_prep[['Tenure', 'MonthlyCharge', 'Techie', 'Phone',
                  'Multiple', 'OnlineSecurity', 'TechSupport', 'StreamingTV',
                  'StreamingMovies', 'PaperlessBilling', 'Marital_Widowed',
                  ↪ 'Gender_Male',
                  'Contract_One year', 'Contract_Two Year',
                  ↪ 'InternetService_Fiber Optic',
                  'InternetService_None', 'PaymentMethod_Electronic Check']]
expTest = sm.add_constant(expTest)

# Data frame of dependent test data
depTest = df_prep['Churn']

yHat = model_reduced.predict(expTest)
prediction = list(map(round, yHat))
```

```
[30]: ## E2 Confusion Matrix and Accuracy Calculation

from sklearn.metrics import confusion_matrix, accuracy_score

# Confusion Matrix
cm = confusion_matrix(depTest, prediction)
print (f'Confusion Matrix: \n {cm}')

# Accuracy Score
print(f'Test accuracy = {accuracy_score(depTest, prediction)}')
```

Confusion Matrix:

```
[[6908  442]
 [ 521 2129]]
```

Test accuracy = 0.9037

4.2.3 E3. Code

See code attached, in D208_PA_MendezD_Task2_Attempt2.ipynb

5 Part V: Data Summary and Implications

5.1 F. Summary of Findings

5.1.1 F1. Results of Data Analysis

- Logistic Regression Equation for the Reduced Model

Let the linear combination of predictor variables and their coefficients, $\mathbf{X}\beta$, be defined by

$$\begin{aligned}\mathbf{X}\beta = & -4.73 - 0.12 \cdot \text{Tenure} + 0.04 \cdot \text{MonthlyCharge} + 1.1 \cdot \text{Techie} \\ & - 0.29 \cdot \text{Phone} + 0.45 \cdot \text{Multiple} - 0.24 \cdot \text{OnlineSecurity} - 0.19 \cdot \text{TechSupport} \\ & + 1.33 \cdot \text{StreamingTV} + 1.51 \cdot \text{StreamingMovies} + 0.16 \cdot \text{PaperlessBilling} + 0.19 \cdot \text{MaritalWidowed} \\ & + 0.28 \cdot \text{GenderMale} - 3.4 \cdot \text{ContractOneyear} - 3.5 \cdot \text{ContractTwoYear} \\ & - 2.14 \cdot \text{InternetServiceFiberOptic} - \text{InternetServiceNone} + 0.48 \cdot \text{PaymentMethodElectronicCheck}\end{aligned}$$

Then the logistic regression equation $\pi(\mathbf{X})$, where π is the probability, for the reduced model is

$$\pi(\mathbf{X}) = \frac{1}{1 + \exp(-\mathbf{X}\beta)}$$

(Logistic Regression, n.d.)

- Interpretation of the Coefficients of the Reduced Model
 - The intercept -4.73 is the baseline probability when all explanatory variables are zero.
 - All else constant, for each additional unit of tenure, the probability of Churn decreases by 0.12.
 - All else constant, for each additional unit of monthly charge, the probability of Churn increases by 0.04.
 - All else constant, being a techie increases the probability of Churn by 1.1 relative to not being a techie.
 - All else constant, having a phone decreases the probability of Churn by 0.29.
 - All else constant, having multiple phone lines increases the probability of Churn by 0.45.
 - All else constant, having online security decreases the probability of Churn by 0.24.
 - All else constant, having tech support decreases the probability of Churn by 0.19.
 - All else constant, having the streaming TV service increases the probability of Churn by 1.33.
 - All else constant, having the streaming movie service increases the probability of Churn by 1.50.
 - All else constant, having paperless billing increases the probability of Churn by 0.16.
 - All else constant, being widowed increases the probability of Churn by 0.19.
 - All else constant, being male increases the probability of Churn by 0.28.
 - All else constant, having a one year contract decreases the probability of Churn by 3.4.
 - All else constant, having a two year contract decreases the probability of Churn by 3.5.
 - All else constant, having fiber optic internet service decreases the probability of Churn by 2.14.
 - All else constant, having no internet service decreases the probability of Churn by 1.
 - All else constant, having electronic check as payment method increases the probability of Churn by 0.48.

- Statistical and Practical Significance of the Reduced Model
 - As shown in E1, with a p-value of $p = 18.9$, at $\alpha = 0.05$, there is evidence to conclude that the model is statistically significant and the additional explanatory variables do not significantly improve the model. However, the model may lack practical significance, as the model does not provide clear insights into why customers churn so it may not be entirely useful in retention strategies.
- Limitations of the Data Analysis
 - A limitation of this data analysis is the assumption that any of these explanatory variables have a causal relationship between them and the dependent variable. Additionally, with 17 explanatory variables, the model could be overfit and result in poor predictive performance.

[31]: *## F1 Regression Equation Coefficients*

```
coefficients = model_reduced.params
print('Coefficients:')
print(coefficients)
```

```
Coefficients:
const                -4.728321
Tenure               -0.115787
MonthlyCharge        0.037535
Techie               1.098711
Phone               -0.289799
Multiple             0.451621
OnlineSecurity      -0.238849
TechSupport         -0.190685
StreamingTV         1.334490
StreamingMovies     1.505083
PaperlessBilling    0.163328
Marital_Widowed    0.192669
Gender_Male         0.275927
Contract_One year   -3.402297
Contract_Two Year   -3.499849
InternetService_Fiber Optic -2.136037
InternetService_None -0.999832
PaymentMethod_Electronic Check 0.475893
dtype: float64
```

5.1.2 F2. Course of Action

Since the Adjusted R^2 of the reduced model was only 0.62, my recommended course of action would be to seek to further improve the model by seeking alternative customer information that might be associated with Churn.

6 Part VI: Demonstration

6.1 G. Panopto Video

<https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=72c62ab6-f7c9-4da4-bf65-b1a6011944d8>

6.2 H. Acknowledgement of Web Sources

Getting started - statsmodels 0.15.0 (+73). (n.d.). <https://www.statsmodels.org/devel/gettingstarted.html>

pandas.get_dummies — pandas 2.2.2 documentation. (n.d.). https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.get_dummies.html

warnings — Warning control — Python 3.10.0 documentation. (n.d.). Docs.python.org. <https://docs.python.org/3/library/warnings.html#warnings.filterwarnings>

Verma, V. (2020, October 24). Feature Selection using Wrapper Method - Python Implementation. Analytics Vidhya. <https://www.analyticsvidhya.com/blog/2020/10/a-comprehensive-guide-to-feature-selection-using-wrapper-methods-in-python/>

Logistic Regression using Statsmodels. (2020, July 17). GeeksforGeeks. <https://www.geeksforgeeks.org/logistic-regression-using-statsmodels/>

Bobbitt, Z. (2021, November 29). How to Perform a Likelihood Ratio Test in Python. Statology. <https://www.statology.org/likelihood-ratio-test-in-python/>

6.3 I. Acknowledgement of Sources

Middleton, K. (2022, November). D208 - Webinar Getting Started with D208 Part I. WGU. <https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=15e09c73-c5aa-439d-852f-af47001b8970>

Middleton, K. (2022, November). D208 - Webinar Getting Started with D208 Part II. WGU. <https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=39bbe2db-de7d-4bf5-913b-af5c0003da9d>

FAQ: What are pseudo R-squareds? UCLA: Statistical Consulting Group. <https://stats.oarc.ucla.edu/other/mult-pkg/faq/general/faq-what-are-pseudo-r-squareds/>

Stoltzfus, J. C. (2011). Logistic Regression: A Brief Primer. Academic Emergency Medicine. <https://pubmed.ncbi.nlm.nih.gov/21996075/>

Logistic Regression (n.d.). Online.stat.psu.edu. <https://online.stat.psu.edu/stat462/node/207/>

likelihood test? <https://www.statology.org/likelihood-ratio-test-in-r/>
<https://sites.warnercnr.colostate.edu/gwhite/wp-content/uploads/sites/73/2017/04/LikelihoodRatioTests.pdf>